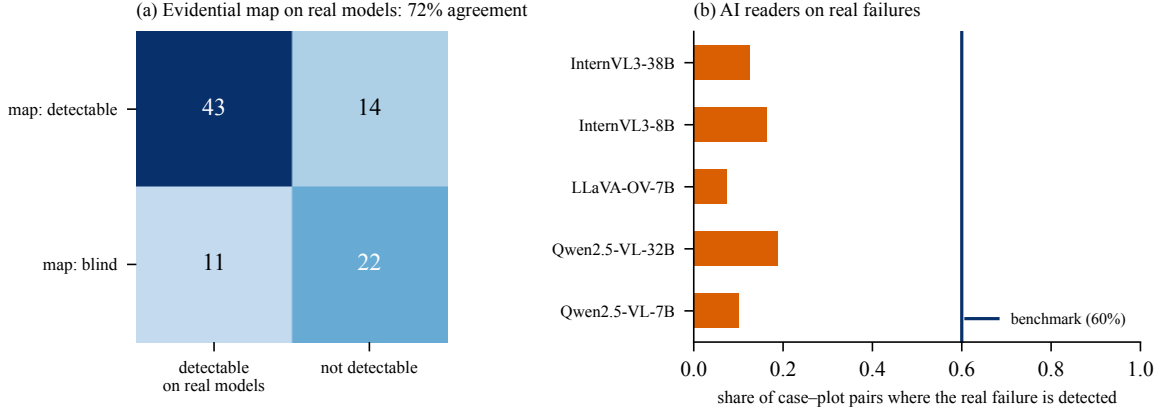
(a) Evidential map on real models: 72% agreement
map: detectable
map: blind
43
14
11
22
detectable
on real models
not detectable
(b) AI readers on real failures
InternVL3-38B
InternVL3-8B
LLaVA-OV-7B
Qwen2.5-VL-32B
Qwen2.5-VL-7B
benchmark (60%)
0.0
0.2
0.4
0.6
0.8
1.0
share of case–plot pairs where the real failure is detected